

Day 2 - Kubernetes - Kube Objects & Namespace

Task: Setup minikube at your local and explore creating namespaces for this class

- So, inside your Kubernetes cluster, you have worker nodes.
- You have Worker Node1 and worker Node2. by default, 2 nodes get created. If you want, you can increase the size of your worker, node, or but by default 2 nodes will be created to form a cluster.
- You need minimum of 2 worker nodes. So, inside your cluster when you use your deployment file. So yesterday we wrote deployment dot yaml file, right? So this is the deployment manifest file. Just written in Yaml format.
- So we use this deployment file to deploy your application inside your cluster.
- So inside your worker node, we have this deployment.
- inside the deployment. You have your pods so you can have multiple pods inside your deployment. You can have N number of Pods. within your pod You have the containers, so your containers will be running inside your pod.
- and your application will be running inside your container. So to access your application. If you want to access your application through the Internet, you cannot directly access it. You have to create an object or resource for these services.
- So you, with the help of services you can, you will be able to access the application that is running inside your Kubernetes cluster. So you have different types of services.
- So you have different types of services like cluster IP. So by default, all your pods, all the resources that is getting created with a cluster IP only.
- So cluster IP is like, which provides you an internal IP [internal IP or internal communication] So you can use this cluster.
- If by default all your containers, will you can communicate with your containers. You can communicate with your pods within the cluster. So for the internal communication we have cluster IP.
- By default, this cluster IP will be attached to your resources that we create inside the cluster. Okay, but with the cluster IP, we can only reach the resources within the cluster. You can access it only within the cluster.

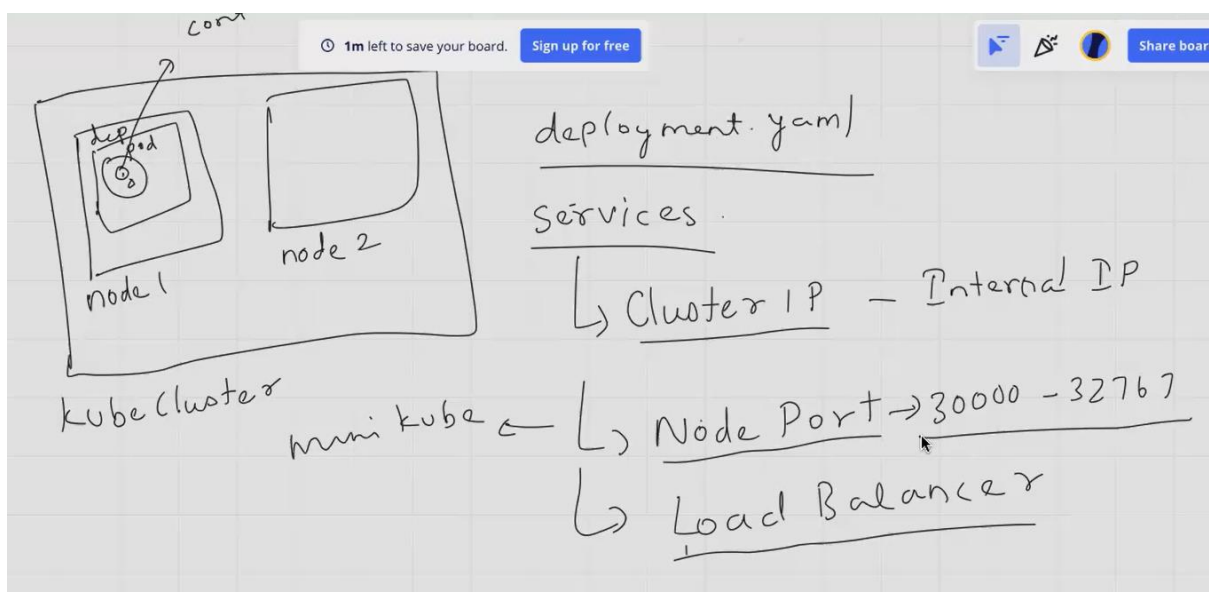
- Outside of the cluster, you cannot access it. so we use a service called node port or if you're running a Kubernetes cluster like how we do it in your cloud environment. We use a service called load balancer.
- So in case if you're running it in on premise, you can either go with the node port. Since we are using aws Cloud Provider, we make use of load balancer service.
- So **this load balancer, and node port will help you to access the application which is running inside the cluster through the Internet.** So via Browser, you'll be able to access the application that is running inside your cluster.
- So typically, we use only load balancer. We don't go with a node port for a local on-premise architecture. You're in your local. If you're setting up Kubernetes cluster, so we will also see how to set up Kubernetes cluster in the local machine using mini kube. In that case you can go with Node port.

- **we have different services.**

1. **cluster IP -internal IP**

2. **Node port** [minikube]. port no between 30,000 to 32,767. you can only configure the port number in this range. Outside of the range you cannot configure. That is the reason we use only node port for local configuration. If you're maintaining a Kubernetes cluster within in the local. then you can go for the in your local machine. Then you can go for node port, because in our laptops and PC. We don't have a load balancer.

3. **Load balancer**



Duplicate the same machine and create a dir.

```
ubuntu@ip-172-31-38-243:~$ mkdir test
ubuntu@ip-172-31-38-243:~$ cd test/
ubuntu@ip-172-31-38-243:~/test$ ls
ubuntu@ip-172-31-38-243:~/test$ clea
```

you can get the Api version for every resource; you have a different Api version for deployment. You have one Api version. If you're not aware of it, initial stages, we will not know which resource is using which Api version. Okay? So, for that, you use a command called this kubeCtl

```
ubuntu@ip-172-31-38-243:~/test$ kubectl api-resources | grep deployment
```

Deployment file [deploy.yaml]

```
ubuntu@ip-172-31-38-243:~/test$ vi deploy.yaml
ubuntu@ip-172-31-38-243:~/test$
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-dep
  labels:
    app: web
spec:
  replicas: 3
  selector:
    matchLabels:
      app: web
  template:
    metadata:
      labels:
        app: web
    spec:
      containers:
        - name: nginx-con
          image: nginx
          ports:
            - containerPort: 80
```

my nginx container will be created inside my deployment.

Code explanation:

This is a Kubernetes YAML file defining a Deployment object. Let's break down the key components:

1. API Version and Kind:

apiVersion: apps/v1: Specifies the API version being used, which is apps/v1 for Deployments.

kind: Deployment: Declares the type of object being defined, which is a Deployment.

2. Metadata:

name: nginx-dep: Sets the name of the Deployment to nginx-dep.

labels: Assigns labels to the Deployment for identification and organization.

app: web: Labels the Deployment as belonging to the web application.

3. Spec:

replicas: 3: Specifies that three replicas of the application should be deployed.

selector: Defines how the Deployment selects which Pods to manage.

match Labels: Selects Pods with the label app: web.

template: Defines the template for creating Pods.

metadata: Metadata for the Pods created from this template.

labels: Assigns labels to the Pods.

app: web: Labels the Pods as belonging to the web application.

spec: Specifies the details of the Pods.

containers: Defines the container(s) to run in the Pods.

- name: nginx-con: Sets the name of the container to nginx-con.

image: nginx: Specifies the Docker image to use for the container, which is the official Nginx image.

ports: Exposes port 80 on the container.

- containerPort: 80: Exposes port 80 within the container.

Overall, this YAML file creates a Deployment named nginx-dep that deploys three replicas of an Nginx web server. The Pods created by this Deployment will be labeled with app: web and will run the Nginx image, exposing port 80

```
ubuntu@ip-172-31-38-243:~$ ls
app1 test
ubuntu@ip-172-31-38-243:~$ cd test/
ubuntu@ip-172-31-38-243:~/test$ ls
deploy.yaml
ubuntu@ip-172-31-38-243:~/test$ kubectl apply -f deploy.yaml
deployment.apps/nginx-dep created
ubuntu@ip-172-31-38-243:~/test$ kubectl get all
NAME                                READY   STATUS              RESTARTS   AGE
pod/nginx-dep-5c979f95d4-fjbhw      0/1     ContainerCreating   0           8s
pod/nginx-dep-5c979f95d4-h8mhl      0/1     ContainerCreating   0           8s
pod/nginx-dep-5c979f95d4-p9tw7      0/1     ContainerCreating   0           8s

NAME                                TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
service/kubernetes                  ClusterIP     10.100.0.1   <none>        443/TCP    12m

NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/nginx-dep           0/3     3            0           8s

NAME                                DESIRED   CURRENT   READY   AGE
replicaset.apps/nginx-dep-5c979f95d4 3         3         0       8s
ubuntu@ip-172-31-38-243:~/test$
```

now we have created the deployment in order to access the deployment through the browser. So now I want to access my nginx application through the Internet. So, what do I need? I need a service resource or a service object.

I have to create a manifest file or a configuration file for creating my service Object/file.

```
ubuntu@ip-172-31-38-243:~/test$ vi service.yaml
ubuntu@ip-172-31-38-243:~/test$ kubectl api-resources | grep deployment
deployments          deploy      apps       true         Deployment
```

```
ubuntu@ip-172-31-38-243:~/test$ kubectl explain deployment
KIND:      Deployment
VERSION:   apps/v1

DESCRIPTION:
  Deployment enables declarative updates for Pods and ReplicaSets.

FIELDS:
  apiVersion    <string>
    APIVersion defines the versioned schema of this representation of an
    object. Servers should convert recognized schemas to the latest internal
    value, and may reject unrecognized values. More info:
    https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#resources
```

```
ubuntu@ip-172-31-38-243:~/test$ kubectl explain services
KIND:      Service
VERSION:   v1
```

if you want to get the Api version of you don't know what is the Api version of service? We are going to create the manifest file. So, we don't know what is the Api version of service. Then you can give Ctl Explain service. It will give you the Api version

service file/object to expose my data via browser:

service.yaml:

```
ubuntu@ip-172-31-38-243:~/test$ vi service.yaml
```

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: web
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: LoadBalancer
```

Target port no means container port no. service runs via load balancer

To check for any error, check in this page

YAML Lint

Paste in your YAML and click "Go" - we'll tell you if it's valid or not, and give you a nice clean UTF-8 version of it.

```

1  ---
2  apiVersion: v1
3  kind: Service
4  metadata:
5    name: my-service
6  spec:
7    selector:
8      app: web
9    ports:
10     - protocol: TCP
11       port: 80
12       targetPort: 8080
13
14
15
16
17
18
19
20

```

Go ☒ Reformat (strips comments) ☒ Resolve aliases

Valid YAML!

```
ubuntu@ip-172-31-38-243:~/test$ kubectl apply -f service.yaml
service/my-service created
```

```
ubuntu@ip-172-31-38-243:~/test$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.100.0.1	<none>	443/TCP	24m
my-service	LoadBalancer	10.100.210.61	a0eab30a4d2fb480eb41d5d709860515-1008349478.ap-south-1.elb.amazonaws.com	80:31260/TCP	17s

```
ubuntu@ip-172-31-38-243:~/test$
```

```
ubuntu@ip-172-31-38-243:~/test$ kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/nginx-dep-5c979f95d4-fjbhw	1/1	Running	0	12m
pod/nginx-dep-5c979f95d4-h8mhl	1/1	Running	0	12m
pod/nginx-dep-5c979f95d4-p9tw7	1/1	Running	0	12m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.100.0.1	<none>	443/TCP	25m
service/my-service	LoadBalancer	10.100.210.61	a0eab30a4d2fb480eb41d5d709860515-1008349478.ap-south-1.elb.amazonaws.com	80:31260/TCP	38s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/nginx-dep	3/3	3	3	12m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/nginx-dep-5c979f95d4	3	3	3	12m

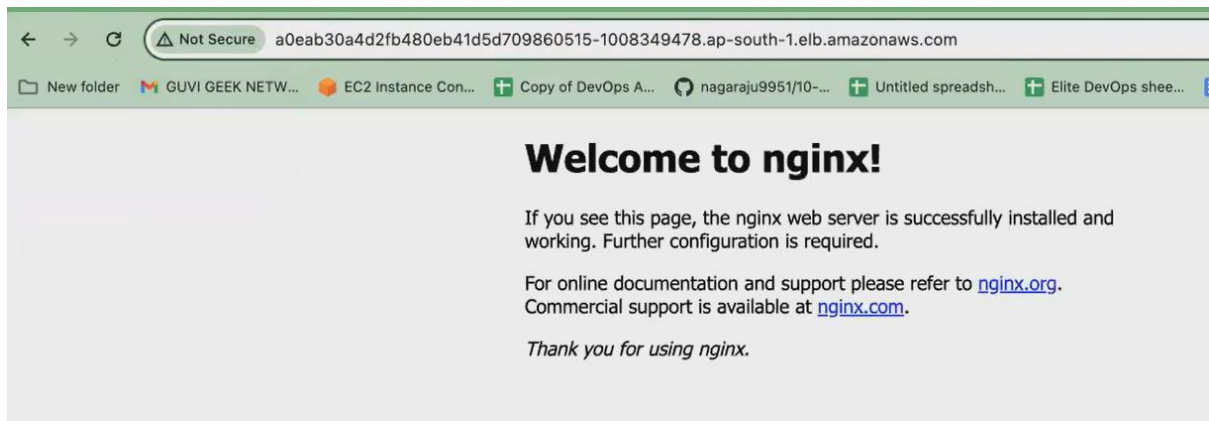
```
ubuntu@ip-172-31-38-243:~/test$
```

Open the port no:80. Copy the external ip and paste it in browser

```
ubuntu@ip-172-31-38-243:~/test$ kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/nginx-dep-5c979f95d4-fjbhw	1/1	Running	0	12m
pod/nginx-dep-5c979f95d4-h8mhl	1/1	Running	0	12m
pod/nginx-dep-5c979f95d4-p9tw7	1/1	Running	0	12m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP
service/kubernetes	ClusterIP	10.100.0.1	<none>
service/my-service	LoadBalancer	10.100.210.61	a0eab30a4d2fb480eb41d5d709860515-1008349478.ap-south-1.elb.amazonaws.com



2 worker node created in same cluster, but running diff ip and machine

Instances (1/5) Info

Find Instance by attribute or tag (case-sensitive) All states

Name	Instance ID	Instance state	Instance type	Status check
guvi-ng-1bb39f85-Node	i-08bc52bbc44b65548	Running	t2.small	2/2 checks passed
eks-demo	i-0b34514b3034b0e6f	Running	t2.medium	2/2 checks passed
guvi-ng-1bb39f85-Node	i-04579bd0e11ef299b	Running	t2.small	2/2 checks passed

i-08bc52bbc44b65548 (guvi-ng-1bb39f85-Node)

Details Status and alarms Monitoring Security Networking Storage Tags

Instance summary Info

Instance ID: i-08bc52bbc44b65548 (guvi-ng-1bb39f85-Node)

Public IPv4 address: 65.0.124.137 | [open address](#)

Private IPv4 addresses: 192.168.39.194, 192.168.54.106

Public IPv4 DNS: ec2-65-0-124-137.ap-south-1.compute.amazonaws.com | [open address](#)

IPv6 address: -

Instance state: Running

NAMESPACE:

- So, suppose assume that you work on a project. You have multiple people working on the same Kubernetes cluster.
- what happened? You and your team member unknowingly you create a deployment with the same name.
- What will happen? You have a conflict there.
- what will happen. You are deleting instead of deleting your deployment, you delete your teammates Deployment. So, what will happen in order to avoid all those things, to maintain the isolation of your resources.

- We go with namespaces so you can create namespaces to isolate the resource creation inside your Kubernetes cluster.
- So, I will tell you how to create the namespace
- so, we can create namespace your any resources, even with the single line command, or you can write a yaml file. So now what we have done. We have written a Yaml file for creating deployment for creating services. You have created yaml file. You can also create the deployment with a single line command also.
- And same way, you can create all the resources, not only deployment, you can create your port, you can create your service object with a single line. We call it as kube ctl cli command.
- Without creating the manifest file. You can go with the kubectl cli Command for emergency purpose.
- You can go with it, but it is always good to follow or create your resources via manifest file only then you will have a proper documentation.

1. Deployment file

2. Service file [used to access the application via browser. That time load balancer created]

✓ Deployment file:

- Create a new dire [test2].
- Mydeployment [deployment file name]
- Httpd [image name]

```
ubuntu@ip-172-31-38-243:~$ cd test2
ubuntu@ip-172-31-38-243:~/test2$ ls
ubuntu@ip-172-31-38-243:~/test2$ kubectl create deployment mydeployment --image=httpd --port=80
deployment.apps/mydeployment created
ubuntu@ip-172-31-38-243:~/test2$
```

```
ubuntu@ip-172-31-38-243:~/test2$ kubectl get all
NAME                                READY    STATUS    RESTARTS   AGE
pod/mydeployment-f65dd7df4-6rk6f    1/1      Running   0           22s
pod/nginx-dep-5c979f95d4-fjbhw      1/1      Running   0           39m
pod/nginx-dep-5c979f95d4-h8mhl      1/1      Running   0           39m
pod/nginx-dep-5c979f95d4-p9tw7      1/1      Running   0           39m

NAME                                TYPE                CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
service/kubernetes                ClusterIP           10.100.0.1      <none>            443/TCP          51m
service/my-service                 LoadBalancer        10.100.210.61   a0eab30a4d2fb480eb41d5d709860515-1008349478.ap-south-1.elb.amazonaws.com  80:31260/TCP     27m

NAME                                READY    UP-TO-DATE    AVAILABLE    AGE
deployment.apps/mydeployment        1/1      1              1            22s
deployment.apps/nginx-dep          3/3      3              3            39m

NAME                                DESIRED    CURRENT    READY    AGE
replicaset.apps/mydeployment-f65dd7df4  1          1          1        22s
replicaset.apps/nginx-dep-5c979f95d4    3          3          3        39m
ubuntu@ip-172-31-38-243:~/test2$
```

```
ubuntu@ip-172-31-38-243:~/test2$ kubectl get all
NAME                                READY    STATUS    RESTARTS   AGE
pod/mydeployment-f65dd7df4-6rk6f    1/1      Running   0           22s
```

- So you have the name of your deployment.
- so the centre one is your replica set Id [f65dd]
- the last 5-digit 5 character is your pod Id. [6rk]

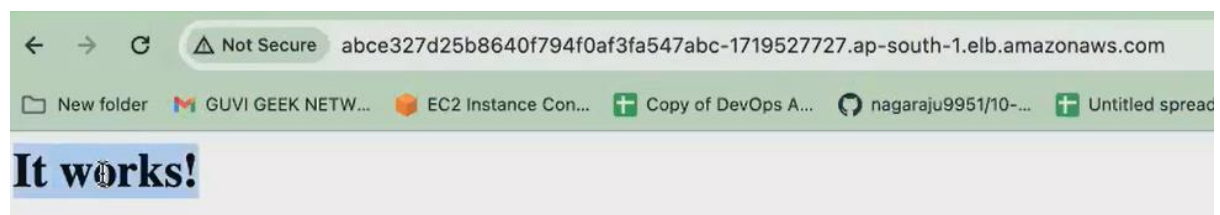
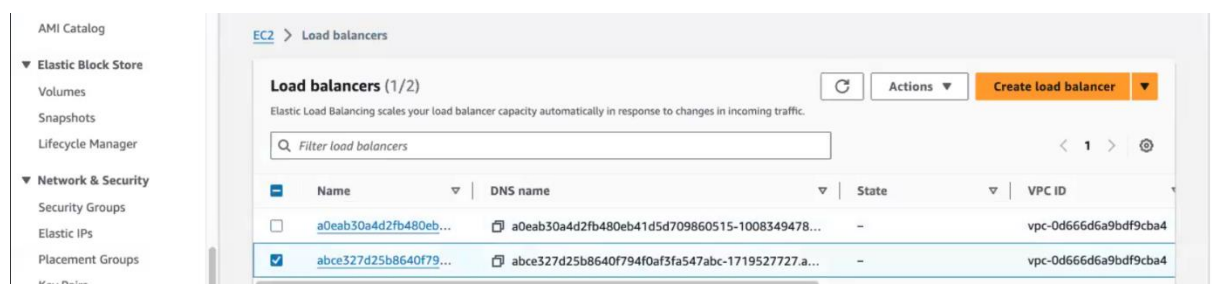
✓ service file:

```
ubuntu@ip-172-31-38-243:~/test2$ kubectl expose deployment mydeployment --type=LoadBalancer --name=apacheservice --port=80 --target-port=80
service/apacheservice exposed
```

```
ubuntu@ip-172-31-38-243:~/test2$ kubectl get svc
NAME                                TYPE                CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
apacheservice                       LoadBalancer        10.100.132.10   abce327d25b8640f794f0af3fa547abc-1719527727.ap-south-1.elb.amazonaws.com  80:30941/TCP     8s
kubernetes                           ClusterIP           10.100.0.1      <none>            443/TCP          54m
my-service                           LoadBalancer        10.100.210.61   a0eab30a4d2fb480eb41d5d709860515-1008349478.ap-south-1.elb.amazonaws.com  80:31260/TCP     29m
```

Copy the ip [abce327d] paste it in browser. Apache server running

Load balancer created:



So, all your resources got created inside the **default namespace**.

```
ubuntu@ip-172-31-38-243:~/test2$ kubectl get ns
NAME                STATUS    AGE
default             Active   57m
kube-node-lease     Active   57m
kube-public         Active   57m
kube-system         Active   57m
```

How to create custom name space:

Kubernetes, everything is created as a resource, everything is treated as objects only. So, your deployment is an object. Your service is an object. And now we are going to create a namespace that is also treated as an object.

Name space file also created via kubectl cli command or manifest file

1. Kubectl cli command:

```
ubuntu@ip-172-31-38-243:~/test2$ kubectl create namespace my-ns
namespace/my-ns created
ubuntu@ip-172-31-38-243:~/test2$ kubectl get ns
NAME                STATUS    AGE
default             Active   61m
kube-node-lease     Active   61m
kube-public         Active   61m
kube-system         Active   61m
my-ns               Active   10s
ubuntu@ip-172-31-38-243:~/test2$
```

2. Manifest file:

```
ubuntu@ip-172-31-38-243:~/test2$ vi namespace.yaml
```

```
apiVersion: v1
kind: Namespace
metadata:
  name: my-ns2
```

Execute the namespace.yaml file:

```
ubuntu@ip-172-31-38-243:~/test2$ kubectl create -f namespace.yaml
namespace/my-ns2 created
ubuntu@ip-172-31-38-243:~/test2$ kubectl get ns
NAME                STATUS    AGE
default             Active    63m
kube-node-lease     Active    63m
kube-public         Active    63m
kube-system         Active    63m
my-ns              Active    2m19s
my-ns2             Active    4s
```

will create a small deployment or pod object inside my namespace.

```
ubuntu@ip-172-31-38-243:~/test2$ vi deployment.yaml
```

Instead of default namespace give our namespace name [ns]

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: apache-deployment
  namespace: my-ns
spec:
  replicas: 3
  selector:
    matchLabels:
      app: apache
  template:
    metadata:
      labels:
        app: apache
    spec:
      containers:
      - name: apache
        image: httpd
        ports:
        - containerPort: 80
```

```
ubuntu@ip-172-31-38-243:~/test2$ vi deployment.yaml
ubuntu@ip-172-31-38-243:~/test2$ kubectl create -f deployment.yaml
deployment.apps/apache-deployment created
```

- what is your deployment name Apache-deployment. It got created inside my custom namespace. So if you want to list it.
- if you give `kubecttl get all` you cannot find this deployment.
- When you give `kubecttl getall` will only display the resources that is there inside your default. Namespace.

```
ubuntu@ip-172-31-38-243:~/test2$ kubecttl get all
NAME                                READY    STATUS    RESTARTS   AGE
pod/mydeployment-f65dd7df4-6rk6f    1/1      Running   0           13m
pod/nginx-dep-5c979f95d4-fjbhw      1/1      Running   0           53m
pod/nginx-dep-5c979f95d4-h8mhl      1/1      Running   0           53m
pod/nginx-dep-5c979f95d4-p9tw7      1/1      Running   0           53m

NAME                                TYPE          CLUSTER-IP    EXTERNAL-IP
service/apacheservice               LoadBalancer  10.100.132.10 abce327d25b8640f794f0af3fa547abc-1719527727.
service/kubernetes                  ClusterIP     10.100.0.1     <none>
service/my-service                   LoadBalancer  10.100.210.61  a0eab30a4d2fb480eb41d5d709860515-1008349478.

NAME                                READY    UP-TO-DATE    AVAILABLE   AGE
deployment.apps/mydeployment        1/1      1              1           13m
deployment.apps/nginx-dep          3/3      3              3           53m

NAME                                DESIRED    CURRENT    READY    AGE
replicaset.apps/mydeployment-f65dd7df4  1          1          1       13m
replicaset.apps/nginx-dep-5c979f95d4    3          3          3       53m
ubuntu@ip-172-31-38-243:~/test2$
```

To list your custom namespaces:

```
ubuntu@ip-172-31-38-243:~/test2$ kubecttl get all -n my-ns
NAME                                READY    STATUS    RESTARTS   AGE
pod/apache-deployment-7749849b9c-2s245 1/1      Running   0           60s
pod/apache-deployment-7749849b9c-jvtdr 1/1      Running   0           60s
pod/apache-deployment-7749849b9c-pfdlm 1/1      Running   0           61s

NAME                                READY    UP-TO-DATE    AVAILABLE   AGE
deployment.apps/apache-deployment        3/3      3              3           61s

NAME                                DESIRED    CURRENT    READY    AGE
replicaset.apps/apache-deployment-7749849b9c  3          3          3       61s
ubuntu@ip-172-31-38-243:~/test2$
```

- we use this namespace just for organizing or for the isolation purpose within. Your Kubernetes cluster mainly for the resource organization and for isolation purpose. You go with your namespace.
- your services, your deployments we go with this namespace. And one more thing is for the access control also.
- Okay, so you can also implement role-based access, which means what you can do. You can further customize this. `You can further control this namespace`. But what will happen? Certain users cannot access your namespaces.
- So, you can either grant a permission to a user or deny a permission to a user to access this namespace. Okay, so there we will be `using a concept called as Role-Based`

Access Control (RBAC). So, we have, like in your Jenkins. You use the concept called RBAC for managing your roles and users.

Commands:

- `kubectl create deployment mydeployment --image=httpd --port=80`
- `kubectl expose deployment mydeployment --type=LoadBalancer --name=apacheservice --port=80 --target-port=80`
- `kubectl create namespace my-ns`

documents link:

<https://docs.google.com/document/d/10dyXzkVgeqWaaJa6AKDaKNfs-tXwu9kgRGwWXm23yz8/edit#heading=h.nf327o3e3njg>

<https://docs.google.com/document/d/10dyXzkVgeqWaaJa6AKDaKNfs-tXwu9kgRGwWXm23yz8/edit?usp=sharing>

minikube official documentation:

<https://minikube.sigs.k8s.io/docs/start/?arch=%2Fwindows%2Fx86-64%2Fstable%2F.exe+download>

Setup minikube at your local: [task]

- ✓ **launch an instance** with t2.medium [its is chargeable.dont use more than 30-40 m].because minikube needs 2 cpu for install.so we go with t2.medium
- ✓ after launch instance, connect and update it. then **install docker**.

✓ **Install kubectl**

✓ **Install minikube**

```
curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64
```

```
sudo install minikube-linux-amd64 /usr/local/bin/minikube && rm minikube-linux-amd64
```

✓ **Minikube start** [minikube start]

✓ **Delete minikube** [minikube delete --all]

minikube is local Kubernetes, focusing on making it easy to learn and develop for Kubernetes.

All you need is [Docker](#) (or similarly compatible) container or a Virtual Machine environment, and Kubernetes is a single command away: `minikube start`

What you'll need

- 2 CPUs or more
- 2GB of free memory
- 20GB of free disk space
- Internet connection
- Container or virtual machine manager, such as: [Docker](#), [QEMU](#), [Hyperkit](#), [Hyper-V](#), [KVM](#), [Parallels](#), [Podman](#), [VirtualBox](#), or [VMware Fusion/Workstation](#)

- binary download If you're using ubuntu you go with Debian package. If you're using centos is on Amazon Linux, you can go with Rpm Base.
- **But binary download is standard for both the distributions.** Binary download You can use it for ubuntu or Amazon Linux or Centos Mission. Better. You go with the binary download itself. Just copy this command, put it in your Ec. 2. Execute it. Once that is done.
- you have to run the Mini cube. Start so Mini Cube will start your cluster.
- Copy and paste the below marked commands in your machine.

1 Installation

Click on the buttons that describe your target platform. For other architectures, see [the release page](#) for a complete list of minikube binaries.

Operating system	<input checked="" type="button" value="Linux"/>	<input type="button" value="macOS"/>	<input type="button" value="Windows"/>		
Architecture	<input checked="" type="button" value="x86-64"/>	<input type="button" value="ARM64"/>	<input type="button" value="ARMv7"/>	<input type="button" value="ppc64"/>	<input type="button" value="S390x"/>
Release type	<input checked="" type="button" value="Stable"/>				
Installer type	<input checked="" type="button" value="Binary download"/>	<input type="button" value="Debian package"/>	<input type="button" value="RPM package"/>		

To install the latest minikube **stable** release on **x86-64 Linux** using **binary download**:

```
curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64
sudo install minikube-linux-amd64 /usr/local/bin/minikube && rm minikube-linux-amd64
```

- Mini Cube start your Mini cube will be started.
- You can see the command control plane running kubectl running kubes running all those things.
- When you give kubeCtl get nodes you can see the nodes running. you will not see multiple nodes running. because kubectl is a single node architecture. You can only see one node. minikube is single node architecture. Used only for development purpose not used for production purpose.
- After install write manifest file
- Don't terminate the instance without delete the cluster. Give command for deleting the cluster [in aws page deleting the cluster is difficult, because del load balancer and all]

```
ubuntu@ip-172-31-38-243:~$ eksctl delete cluster guvi --region ap-south-1
2024-07-19 12:38:49 [i] deleting EKS cluster "guvi"
2024-07-19 12:38:49 [i] will drain 0 unmanaged nodegroup(s) in cluster "guvi"
2024-07-19 12:38:49 [i] starting parallel draining, max in-flight of 1
2024-07-19 12:38:49 [*] failed to acquire semaphore while waiting for all routines to finish: context canceled
2024-07-19 12:38:49 [i] deleted 0 Fargate profile(s)
2024-07-19 12:38:49 [✓] kubeconfig has been updated
2024-07-19 12:38:49 [i] cleaning up AWS load balancers created by Kubernetes objects of Kind Service or Ingress
```

2 Start your cluster

From a terminal with administrator access (but not logged in as root), run:

```
minikube start
```

If minikube fails to start, see the [drivers page](#) for help setting up a compatible container or virtual-machine manager.

Day 3 - Kubernetes - Pods & Services . doubt session

<https://docs.google.com/document/d/1Is4h94KVFliaNxSuBZX9Spui98CYiFOlawjLfYa3wBU/edit>

this doc shared by Naveen mentor

